

Introduction to Clustering

Lieven Clement

Table of contents

1	Introduction	2
1.1	Objective	2
1.2	Example 1	2
1.3	Example 2	2
2	Partition Based Cluster Analysis	2
2.1	K-means Methods	3
2.2	Example	3
3	Hierarchical Cluster Analysis	6
3.1	General Algorithm of Agglomerative Hierarchical Clustering	6
3.2	Intercluster Dissimilarities	6
3.2.1	Single Linkage = Nearest Neighbour	6
3.2.2	Complete Linkage = Furthest Neighbour	7
3.2.3	Average Linkage = Group Average	8
3.3	Cluster Tree	9
4	Toy example	9
4.1	Single linkage	10
4.2	Complete linkage	11
4.3	Average linkage	12
4.4	Example	13
5	Model-based clustering	15
	Session info	15

1 Introduction

1.1 Objective

Objective: grouping of observations into **clusters**, so that

- similar observations appear in the same cluster
- dissimilar observations appear in distinct clusters

→ need for a measure for **similarity** and **dissimilarity**?

1.2 Example 1

Single cell transcriptomics: $n \times p$ Matrix for which

- every column contains the expression levels of one of p genes for n cells
- every row contains the expression levels of p genes for one cell (**sample**)
- Research question: look for groups of cells that have similar gene expression patterns
- Or, look for groups of genes that have similar expression levels across the different cells. This can help us in understanding the regulation and functionality of the genes.

→ both **observations** (rows) and **variables** (columns) can be clustered

1.3 Example 2

Abundance studies: the abundances of n plant species are counted on p plots (habitats)

- look for groups that contain species that live in the same habitats, or, look for groups of habitats that have similar species communities

→ both **observations** (rows) and **variables** (columns) can be clustered

2 Partition Based Cluster Analysis

- Partition based cluster methods require the number of clusters (k) to be specified prior to the start of the algorithm.

2.1 K-means Methods

- To use the k-means clustering algorithm we have to pre-define k , the number of clusters we want to define.
- The k-means algorithm is iterative.
- The algorithm starts by defining k cluster centers (centroids).
- Then the algorithm proceeds as follows
 1. First each observation is assigned to the cluster with the closest center to that observation.
 2. Then the k centers are redefined using the observations in each cluster, i.e. the multivariate means (column means) of all observations in a cluster are used to define each new cluster center.
 3. We repeat these two steps until the centers converge.

2.2 Example


```
library(tidyverse)
data(diabetes, package = "mclust")
class <- diabetes$class
table(class)
```

```
#> class
#> Chemical    Normal    Overt
#>           36         76         33
```

```
head(diabetes)
```

```
#>   class glucose insulin sspg
#> 1 Normal      80      356  124
#> 2 Normal      97      289  117
#> 3 Normal     105      319  143
#> 4 Normal      90      356  199
#> 5 Normal      90      323  240
#> 6 Normal      86      381  157
```

```
mclust::clPairs(diabetes[, -1], diabetes$class)
```



```
diabetesKmeans <- kmeans(diabetes[,-1], centers = 3)
diabetesKmeans
```

```
#> K-means clustering with 3 clusters of sizes 26, 94, 25
#>
#> Cluster means:
#>   glucose   insulin    sspg
#> 1 241.65385 1152.8846  75.69231
#> 2  93.39362  375.5213 166.17021
#> 3 105.04000  525.6000 375.96000
#>
#> Clustering vector:
#>   1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
#>   2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2
#>  21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
#>   2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2
#>  41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60
#>   2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2
#>  61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80
#>   2  2  2  2  2  2  2  2  3  2  3  2  2  2  2  2  2  2  2  2
#>  81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100
#>   2  3  2  2  2  3  3  2  3  3  3  3  3  3  3  2  3  3  3  3
#> 101 102 103 104 105 106 107 108 109 110 111 112 113 114 115 116 117 118 119 120
```

```



#>   3   3   2   2   2   3   2   3   2   2   3   2   1   1   3   1   1   1   1   1
#> 121 122 123 124 125 126 127 128 129 130 131 132 133 134 135 136 137 138 139 140
#>   1   1   1   2   1   1   1   1   1   1   1   3   1   1   2   2   3   3   1   1   1
#> 141 142 143 144 145
#>   1   1   1   1   1
#>
#> Within cluster sum of squares by cluster:
#> [1] 1738796.1 947827.2 687281.9
#> (between_SS / total_SS = 80.6 %)
#>
#> Available components:
#>
#> [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
#> [6] "betweenss"    "size"         "iter"         "ifault"

```

```

mclust::clPairs(diabetes[, -1], diabetes$class, main = "Real"); mclust::clPairs(diabetes[, -1]

```



3 Hierarchical Cluster Analysis

- Distinction between agglomerative and divisive methods
- Agglomerative start from the situation where each individual observations forms its own cluster (so it starts with n clusters). In the next steps clusters are sequentially merged, until finally there is only one cluster with n observations.
- Divisive methods work just the other way around.
- The solution of an hierarchical clustering is thus a sequence of n nested cluster solutions.

3.1 General Algorithm of Agglomerative Hierarchical Clustering

- In step 0 each observations is considered as a cluster (i.e. n clusters).
- Every next step consists of:
 1. merge the two clusters with the smallest intercluster dissimilarity
 2. recalculate the intercluster dissimilarities

In step 0 the intercluster dissimilarity coincides with the dissimilarity between the corresponding observations

→ intercluster dissimilarity?

3.2 Intercluster Dissimilarities


- Represent clusters (e.g. C_1 and C_2) as sets of points $\mathbf{x}_i$ which belong to that cluster
- $d(C_1, C_2)$: intercluster dissimilarity between

We consider three intercluster dissimilarities.

3.2.1 Single Linkage = Nearest Neighbour

$$d(C_1, C_2) = \min_{\mathbf{x}_1 \in C_1; \mathbf{x}_2 \in C_2} d(\mathbf{x}_1, \mathbf{x}_2),$$


i.e. the dissimilarity between C_1 and C_2 is determined by the smallest dissimilarity between a point of C_1 and a point of C_2 .



3.2.2 Complete Linkage = Furthest Neighbour

$$d(C_1, C_2) = \max_{\mathbf{x}_1 \in C_1; \mathbf{x}_2 \in C_2} d(\mathbf{x}_1, \mathbf{x}_2),$$


i.e. the dissimilarity between C_1 and C_2 is determined by the largest dissimilarity between a point of C_1 and a point of C_2 .



3.2.3 Average Linkage = Group Average

$$d(C_1, C_2) = \frac{1}{|C_1||C_2|} \sum_{\mathbf{x}_1 \in C_1; \mathbf{x}_2 \in C_2} d(\mathbf{x}_1, \mathbf{x}_2),$$

i.e. the dissimilarity between C_1 and C_2 is determined by the average dissimilarity between all points of C_1 and all points of C_2 .



3.3 Cluster Tree


Hierarchical nature of the algorithm:

- Nested sequence of clusters $\rightarrow$ visualisation via a tree
- Height of branches indicate the intercluster dissimilarity at which clusters are merged.
- Can used as instrument for deciding the number of clusters in the data

4 Toy example

X1	X2	label
1.50	2.40	1
2.00	2.50	2
2.50	2.25	3
2.00	3.00	4
2.25	3.20	5

```
toy |>
  ggplot(aes(X1, X2, label = label)) +
  geom_point() +
  geom_text(nudge_x = .05)
```




```
toy[,1:2] |> dist()
```

```
#>           1           2           3           4
#> 2 0.5099020
#> 3 1.0111874 0.5590170
#> 4 0.7810250 0.5000000 0.9013878
#> 5 1.0965856 0.7433034 0.9823441 0.3201562
```

4.1 Single linkage

```
toyDist <- toy[,1:2] |> dist()
toySingle <- hclust(toyDist, method = "single")
par(mfrow=c(1,2),pty="s")
plot(X2 ~ X1, toy, xlim = c(1.25,2.75),ylim = c(2,3.5))
```

```
text(toy$X1*1.05,toy$X2,label=toy$label)
plot(toySingle, main = "Single")
```




```
toyDist
```

```
#>      1      2      3      4
#> 2 0.5099020
#> 3 1.0111874 0.5590170
#> 4 0.7810250 0.5000000 0.9013878
#> 5 1.0965856 0.7433034 0.9823441 0.3201562
```

4.2 Complete linkage

```
toyComplete <- hclust(toyDist, method = "complete")
par(mfrow=c(1,2),pty="s")
plot(X2 ~ X1, toy, xlim = c(1.25,2.75),ylim = c(2,3.5))
text(toy$X1*1.05,toy$X2,label=toy$label)
plot(toyComplete, main = "Complete")
```




```
toyDist
```

```
#>           1           2           3           4
#> 2 0.5099020
#> 3 1.0111874 0.5590170
#> 4 0.7810250 0.5000000 0.9013878
#> 5 1.0965856 0.7433034 0.9823441 0.3201562
```

4.3 Average linkage

```
toyAvg <- hclust(toyDist, method = "average")
par(mfrow=c(1,2),pty="s")
plot(X2 ~ X1, toy, xlim = c(1.25,2.75),ylim = c(2,3.5))
text(toy$X1*1.05,toy$X2,label=toy$label)
plot(toyAvg, main = "Average")
```




```
toyDist
```

```
#>      1      2      3      4
#> 2 0.5099020
#> 3 1.0111874 0.5590170
#> 4 0.7810250 0.5000000 0.9013878
#> 5 1.0965856 0.7433034 0.9823441 0.3201562
```

4.4 Example

```
diabetesDist <- dist(diabetes[, -1])
diabetesSingle <- hclust(diabetesDist, method = "single")
plot(diabetesSingle, labels = as.double(diabetes$class), main="single")
```


single



diabetesDist
hclust (*, "single")


```
diabetesComplete <- hclust(diabetesDist, method = "complete")
plot(diabetesComplete, labels = as.double(diabetes$class), main="complete")
```

complete



diabetesDist
hclust (*, "complete")

```
diabetesAverage <- hclust(diabetesDist, method = "average")
plot(diabetesAverage, labels = as.double(diabetes$class), main = "average", cex=0.5)
```



5 Model-based clustering

- [Paper: Fraley and Raftery \(1998\). How Many Clusters? Which Clustering Method? Answers Via Model-Based Cluster Analysis. The Computer Journal, \(41\)8:578-588.](#)
- [EM algorithm \[PDF\]](#)
- Example: see tutorial session

Session info

Session info

```
#> [1] "2026-09-15 16:45:21 CEST"
```

```

#> - Session info -----
#> setting value
#> version R version 4.6.1 (2026-06-24)
#> os      macOS Tahoe 26.1
#> system  aarch64, darwin23
#> ui      X11
#> language (EN)
#> collate nl_BE.UTF-8
#> ctype   nl_BE.UTF-8
#> tz      Europe/Brussels
#> date    2026-09-15
#> pandoc  3.10.2 @ /opt/homebrew/bin/ (via rmarkdown)
#> quarto  1.9.37 @ /usr/local/bin/quarto
#>
#> - Packages -----
#> package      * version date (UTC) lib source
#> cli           3.6.6   2026-04-09 [1] CRAN (R 4.6.0)
#> digest        0.6.39  2025-11-19 [1] CRAN (R 4.6.0)
#> dplyr         * 1.2.1   2026-04-03 [1] CRAN (R 4.6.0)
#> evaluate      1.0.5   2025-08-27 [1] CRAN (R 4.6.0)
#> farver        2.1.2   2024-05-13 [1] CRAN (R 4.6.0)
#> fastmap       1.2.0   2024-05-15 [1] CRAN (R 4.6.0)
#> forcats       * 1.0.1   2025-09-25 [1] CRAN (R 4.6.0)
#> generics      0.1.4   2025-05-09 [1] CRAN (R 4.6.0)
#> ggplot2       * 4.0.3   2026-04-22 [1] CRAN (R 4.6.0)
#> glue          1.8.1   2026-04-17 [1] CRAN (R 4.6.0)
#> gtable        0.3.6   2024-10-25 [1] CRAN (R 4.6.0)
#> hms           1.1.4   2025-10-17 [1] CRAN (R 4.6.0)
#> htmltools     0.5.9   2025-12-04 [1] CRAN (R 4.6.0)
#> jsonlite      2.0.0   2025-03-27 [1] CRAN (R 4.6.0)
#> knitr         1.51    2025-12-20 [1] CRAN (R 4.6.0)
#> labeling      0.4.3   2023-08-29 [1] CRAN (R 4.6.0)
#> lifecycle     1.0.5   2026-01-08 [1] CRAN (R 4.6.0)
#> lubridate     * 1.9.5   2026-02-04 [1] CRAN (R 4.6.0)
#> magrittr      2.0.5   2026-04-04 [1] CRAN (R 4.6.0)
#> mclust        6.1.3   2026-07-05 [1] CRAN (R 4.6.1)
#> ote           0.2.0   2025-08-29 [1] CRAN (R 4.6.0)
#> pillar        1.11.1  2025-09-17 [1] CRAN (R 4.6.0)
#> pkgconfig     2.0.3   2019-09-22 [1] CRAN (R 4.6.0)
#> purrr         * 1.2.2   2026-04-10 [1] CRAN (R 4.6.0)
#> R6            2.6.1   2025-02-15 [1] CRAN (R 4.6.0)
#> RColorBrewer  1.1-3   2022-04-03 [1] CRAN (R 4.6.0)
#> readr         * 2.2.0   2026-02-19 [1] CRAN (R 4.6.0)

```

```

#> rlang          1.2.0    2026-04-06 [1] CRAN (R 4.6.0)
#> rmarkdown      2.31     2026-03-26 [1] CRAN (R 4.6.0)
#> S7              0.2.2    2026-04-22 [1] CRAN (R 4.6.0)
#> scales         1.4.0     2025-04-24 [1] CRAN (R 4.6.0)
#> sessioninfo    1.2.4     2026-06-04 [1] CRAN (R 4.6.0)
#> stringi        1.8.7     2025-03-27 [1] CRAN (R 4.6.0)
#> stringr        * 1.6.0    2025-11-04 [1] CRAN (R 4.6.0)
#> tibble         * 3.3.1    2026-01-11 [1] CRAN (R 4.6.0)
#> tidyr          * 1.3.2    2025-12-19 [1] CRAN (R 4.6.0)
#> tidyselect     1.2.1     2024-03-11 [1] CRAN (R 4.6.0)
#> tidyverse      * 2.0.0    2023-02-22 [1] CRAN (R 4.6.0)
#> timechange     0.4.0     2026-01-29 [1] CRAN (R 4.6.0)
#> tzdb           0.5.0     2025-03-15 [1] CRAN (R 4.6.0)
#> vctrs          0.7.3     2026-04-11 [1] CRAN (R 4.6.0)
#> withr          3.0.3     2026-06-19 [1] CRAN (R 4.6.0)
#> xfun           0.59      2026-06-19 [1] CRAN (R 4.6.0)
#> yaml           2.3.12    2025-12-10 [1] CRAN (R 4.6.0)
#>
#> [1] /Library/Frameworks/R.framework/Versions/4.6/Resources/library
#> * -- Packages attached to the search path.
#>
#> -----

```